

Apache Spark Configuration Cheat Sheet

Scenario 1: Slow Join Job

Problem: Join operation is slow with high shuffle spill.

Solutions:

- Increase shuffle partitions: `spark.sql.shuffle.partitions = 800`
- Optimize executor resources: `spark.executor.memory = 6g`, `spark.executor.cores = 4`
- Cache reused DataFrames to avoid recomputation
- Use broadcast join if one dataset is small:
`df1.join(broadcast(df2), "id")`

Scenario 2: OutOfMemoryError in Executors

Problem: Executors fail due to insufficient memory.

Solutions:

- Increase executor memory: `spark.executor.memory = 8g`
- Use G1GC collector:
`spark.executor.extraJavaOptions = "-XX:+UseG1GC"`
- Avoid large `collect()` or `toPandas()` calls on big datasets

Scenario 3: Skewed Tasks (Stragglers)

Problem: Some tasks take significantly longer due to data skew.

Solutions:

- Enable speculative execution: `spark.speculation = true`
- Use salting techniques for skewed keys
- Repartition skewed data: `df.repartition("key_column")`

Scenario 4: Job Crashes on Full Data Load

Problem: Job works on sample but fails with full dataset.

Apache Spark Configuration Cheat Sheet

Solutions:

- Tune based on sample profiling using Spark UI
- Scale executor memory/cores or cluster size
- Tune shuffle partitions and default parallelism:
spark.sql.shuffle.partitions = 1000
spark.default.parallelism = 800

Scenario 5: High Executor Idle Time

Problem: Executors are idle, causing poor resource usage.

Solutions:

- Enable dynamic allocation:
spark.dynamicAllocation.enabled=true
spark.dynamicAllocation.minExecutors=4
spark.dynamicAllocation.maxExecutors=100
- Reduce shuffle partitions if over-partitioned